



PyTorch Zero To All 9-10


# Sofmax Classifier & Basic CNN

최건우




# Softmax Classifier

In Logistic Regression





# Softmax Classifier





# Softmax Classifier

Softmax function

$$\text{softmax}(z) = \left[ \frac{e^{z_1}}{\sum_{j=1}^3 e^{z_j}}, \frac{e^{z_2}}{\sum_{j=1}^3 e^{z_j}}, \frac{e^{z_3}}{\sum_{j=1}^3 e^{z_j}} \right]$$






# Softmax Classifier

Softmax function

$$\text{softmax}(z) = \left[ \frac{e^{z_1}}{\sum_{j=1}^3 e^{z_j}} \quad \frac{e^{z_2}}{\sum_{j=1}^3 e^{z_j}} \quad \frac{e^{z_3}}{\sum_{j=1}^3 e^{z_j}} \right]$$



# Softmax Classifier


Why 'e'?

$$\text{softmax}(z) = \left[ \frac{e^{z_1}}{\sum_{j=1}^3 e^{z_j}} \quad \frac{e^{z_2}}{\sum_{j=1}^3 e^{z_j}} \quad \frac{e^{z_3}}{\sum_{j=1}^3 e^{z_j}} \right]$$

1.

$$\frac{d}{dx} e^x = e^x$$

2.




[https://www.youtube.com/watch?v=K7HTd\\_Zgr3w&t=329s](https://www.youtube.com/watch?v=K7HTd_Zgr3w&t=329s)





# Softmax Classifier

Cost function



$$cost(W) = - \sum_{j=1}^k y_j \log(p_j)$$

$$cost(W) = -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k y_j^{(i)} \log(p_j^{(i)})$$



# Softmax Classifier

Logistic Regression Cost function

$$\text{cost}(W) = -(y \log H(X) + (1 - y) \log(1 - H(X)))$$

$$\begin{array}{ll} y \rightarrow y_1 & (1-y) \rightarrow y_2 \\ H(x) \rightarrow p_1 & (1-H(x)) \rightarrow p_2 \end{array}$$

$$-(y_1 \log(p_1) + y_2 \log(p_2))$$

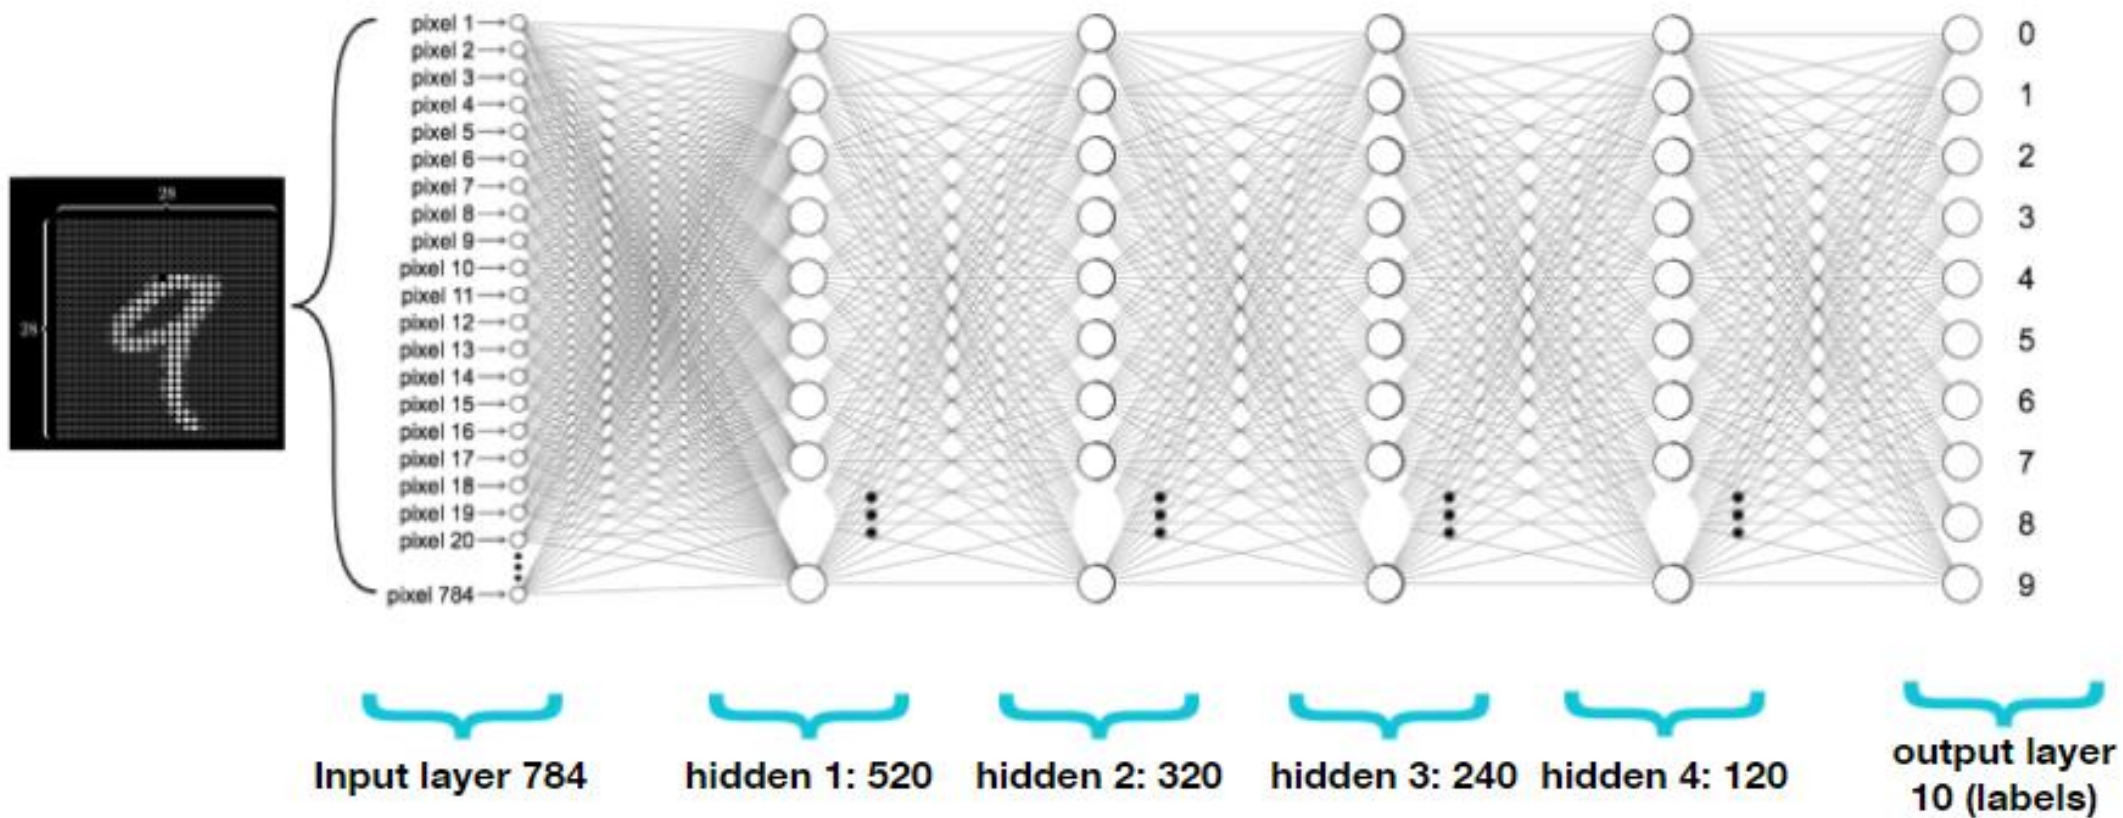
$$-\left(\sum_{i=1}^2 y_i \log p_i\right)$$


Softmax Regression Cost function

$$\text{cost}(W) = -\sum_{j=1}^k y_j \log(p_j)$$



# Softmax Classifier





# Softmax Classifier

```
class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.l1 = nn.Linear(784, 520)
        self.l2 = nn.Linear(520, 320)
        self.l3 = nn.Linear(320, 240)
        self.l4 = nn.Linear(240, 120)
        self.l5 = nn.Linear(120, 10)

    def forward(self, x):
        x = x.view(-1, 784) # Flatten the data (n, 1, 28, 28) -> (n, 784)
        x = F.relu(self.l1(x))
        x = F.relu(self.l2(x))
        x = F.relu(self.l3(x))
        x = F.relu(self.l4(x))
        return self.l5(x)
```

```
criterion = nn.CrossEntropyLoss()
```

```
...
for batch_idx, (data, target) in enumerate(train_loader):
    data, target = Variable(data), Variable(target)
    optimizer.zero_grad()
    output = model(data)
    loss = criterion(output, target)
    loss.backward()
    optimizer.step()
```





기획  
공방

# Softmax Classifier

```
def train(epoch):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        if batch_idx % 10 == 0:
            print('Train Epoch: {} | Batch Status: {}/{} ({:.0f}%) | Loss: {:.6f}'.format(
                epoch, batch_idx + len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))


def test():
    model.eval()
    test_loss = 0
    correct = 0
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        output = model(data)
        # sum up batch loss
        test_loss += criterion(output, target).item()
        # get the index of the max
        pred = output.data.max(1, keepdim=True)[1]
        correct += pred.eq(target.data.view_as(pred)).cpu().sum()

    test_loss /= len(test_loader.dataset)
    print(f'=====\nTest set: Average loss: {test_loss:.4f}, Accuracy: {correct}/{len(test_loader.dataset)}\n({100. * correct / len(test_loader.dataset):.0f}%)')

if __name__ == '__main__':
    since = time.time()
    for epoch in range(1, 10):
        epoch_start = time.time()
        train(epoch)
        m, s = divmod(time.time() - epoch_start, 60)
        print(f'Training time: {m:.0f}m {s:.0f}s')
        test()
        m, s = divmod(time.time() - epoch_start, 60)
        print(f'Testing time: {m:.0f}m {s:.0f}s')

    m, s = divmod(time.time() - since, 60)
    print(f'Total Time: {m:.0f}m {s:.0f}s\nModel was trained on {device}!')
```

Test set: Average loss: 0.0016, Accuracy: 9691/10000 (97%)  
Testing time: 0m 22s  
Total Time: 3m 39s  
Model was trained on cpu!



# Softmax Classifier

```
class Net(nn.Module):
```

```
    def __init__(self):
```

```
        super(Net, self).__init__()
```

```
        self.l1 = nn.Linear(784, 650)
```

```
        self.l2 = nn.Linear(650, 530)
```

```
        self.l3 = nn.Linear(530, 420)
```

```
        self.l4 = nn.Linear(420, 350)
```

```
        self.l5 = nn.Linear(350, 290)
```

```
        self.l6 = nn.Linear(290, 200)
```

```
        self.l7 = nn.Linear(200, 150)
```

```
        self.l8 = nn.Linear(150, 80)
```

```
        self.l9 = nn.Linear(80, 10)
```

```
    def forward(self, x):
```

```
        x = x.view(-1, 784) # Flatten the data (n, 1, 28, 28) -> (n, 784)
```

```
        x = F.relu(self.l1(x))
```

```
        x = F.relu(self.l2(x))
```

```
        x = F.relu(self.l3(x))
```

```
        x = F.relu(self.l4(x))
```

```
        x = F.relu(self.l5(x))
```

```
        x = F.relu(self.l6(x))
```

```
        x = F.relu(self.l7(x))
```

```
        x = F.relu(self.l8(x))
```

```
        return self.l9(x)
```

Test set: Average loss: 0.0361, Accuracy: 1135/10000 (11%)  
Testing time: 0m 30s  
Total Time: 5m 17s  
Model was trained on cpu!

```
class Net(nn.Module):
```

```
    def __init__(self):
```

```
        super(Net, self).__init__()
```

```
        self.l1 = nn.Linear(784, 420)
```

```
        self.l2 = nn.Linear(420, 120)
```

```
        self.l3 = nn.Linear(120, 10)
```

```
    def forward(self, x):
```

```
        x = x.view(-1, 784) # Flatten the data (n, 1, 28, 28) -> (n, 784)
```

```
        x = F.relu(self.l1(x))
```

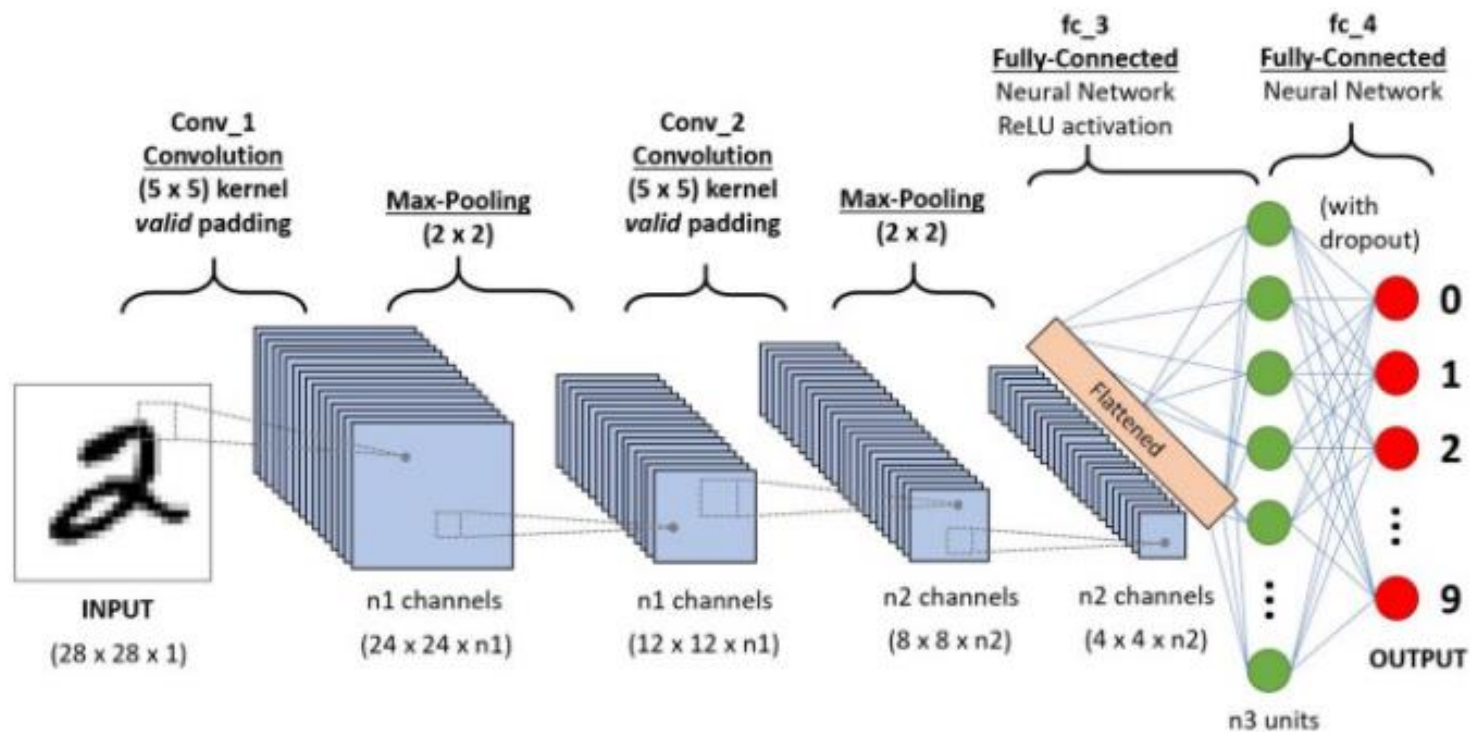
```
        x = F.relu(self.l2(x))
```

```
        return self.l3(x)
```

Test set: Average loss: 0.0022, Accuracy: 9574/10000 (96%)  
Testing time: 0m 28s  
Total Time: 2m 18s  
Model was trained on cpu!



# Basic CNN





# Basic CNN

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image


4		

Convolved  
Feature






# Basic CNN



Stride 1 with Padding



Feature Map



# Basic CNN

Single depth slice


1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters  
and stride 2



6	8
3	4





# Basic CNN

```
class Net(nn.Module):  
  
    def __init__(self):  
        super(Net, self).__init__()  
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)  
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)  
        self.mp = nn.MaxPool2d(2)  
        self.fc = nn.Linear(320, 10)  
  
    def forward(self, x):  
        in_size = x.size(0)  
        x = F.relu(self.mp(self.conv1(x)))  
        x = F.relu(self.mp(self.conv2(x)))  
        x = x.view(in_size, -1) # flatten the tensor  
        x = self.fc(x)  
        return F.log_softmax(x)
```

input: 1x28x28

10x24x24 -> 10x12x12

20x8x8 -> 20x4x4=320



# Basic CNN

```
class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(???, 10)

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```


**RuntimeError:** mat1 and mat2 shapes cannot be multiplied (64x320 and 100x10)

# Basic CNN

```
class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(320, 10)

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```



$$\text{MaxPool}(\text{Relu}(x)) = \text{Relu}(\text{MaxPool}(x))$$





# Basic CNN

```
model = Net()

optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)

def train(epoch):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = Variable(data), Variable(target)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)
        loss.backward()
        optimizer.step()
        if batch_idx % 10 == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))

def test():
    model.eval()
    test_loss = 0
    correct = 0
    for data, target in test_loader:
        data, target = Variable(data, volatile=True), Variable(target)
        output = model(data)
        # sum up batch loss
        test_loss += F.nll_loss(output, target, size_average=False).data
        # get the index of the max log-probability
        pred = output.data.max(1, keepdim=True)[1]
        correct += pred.eq(target.data.view_as(pred)).cpu().sum()

    test_loss /= len(test_loader.dataset)
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)\n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))

for epoch in range(1, 10):
    train(epoch)
    test()
```

Test set: Average loss: 0.0576, Accuracy: 9814/10000 (98%)



D-ai-ving

**감사합니다!**